

Examining readability of text generated through GPT-3

Aron Winkler

University of Tübingen

winkler.aron5@gmail.com

Abstract

This paper examines readability-related metrics of AI-generated text, specifically texts produced by OpenAI's GPT-3. Generation was performed across 10 topics and the 6 proficiency levels defined by the CEFR framework, and outputs were compared with Simple Wikipedia and (standard) Wikipedia articles on the same topics.

1 Introduction

AI text generation has experienced a massive rise in popularity with the media reach achieved by the release of OpenAI's GPT-3 (Brown et al., 2020) and its chat GUI interface ChatGPT¹. While early public opinion focused on how - in the school setting - it would be students who would benefit from the new system by having their schoolwork done for them, little was said about how AI text generation could become an asset for didacticians in generating teaching material.

AI text generation might become an important tool for teachers, allowing them to generate learning material about topics and events that engage their students that might not be available on the market. An interesting setting for such a phenomenon is language learning, where it is necessary to produce learning material for students at different proficiency levels: whether to create ex-novo or simplify existing text material, AI text-generation tools can easily be imagined as having an enabling effect.

It is however not obvious that AI generation tools reach the desired quality or adherence to the didactic goals that teachers set out. Owing to the increased popularity of these tools, more research is desirable on their applicability in a number of different settings.

This work aims to make a contribution to those questions by verifying the readability level of text

content generated by GPT-3 with regard to 10 topics and the 6 proficiency levels defined by the CEFR framework. A number of complexity metrics are computed on the generated texts and compared with Simple Wikipedia² as well as Wikipedia³ articles on the same topics in order to gauge the true readability levels of the model outputs.

2 Dataset and metrics

2.1 Dataset

OpenAI's REST api was employed to produce texts with the "text-davinci-003" model (a GPT-3 variant) across 10 topics⁴ and 6 CEFR proficiency levels⁵. Topics were selected without any particular guiding principle other than a general goal of variety between events, entities, and phenomena. The prompts fed to the model followed the template:

Write a text for learners of English at the [level] level about [topic].

In addition to the 60 texts obtained in this way, the dataset was enriched by Simple Wikipedia and Wikipedia articles corresponding to each of the 10 topics, resulting in an additional 20 dataset entries.

2.2 Preprocessing

Preprocessing of the dataset comprised enrichment with POS and dependency annotation. Part-of-speech tagging was achieved with nltk (Loper and Bird, 2002) and dependency relations were added through nltk in conjunction with CoreNLP models (Manning et al., 2014).

¹<https://openai.com/blog/chatgpt>

Metric	Abbreviation
Sentence length	ASL
Noun to verb ratio	NVR
Type Token Ratio (first 100 tokens)	TTR
Clauses Per Sentence	CPS
Length of Longest Dependency Link	LLDL
Parse Tree Depth	PTD
Ratio of words in top 3000 English words	WIT

Table 1: Readability metrics

2.3 Metrics

Table 1 lists the readability metrics used in this experiment. This specific selection was inspired by (Venturi et al., 2015), but many more examples of their usage can be found.

This feature set almost exclusively targets lexis and syntax, with little regard for morphology outside of the noun-to-verb ratio feature. While this is a limitation of this set, the language of this experiment is solely English, a language not necessarily known for a complex morphological system.

The feature "Words in Top 3000 English words" uses unigram frequencies from the Google Web Trillion Word corpus⁶ to isolate the 3000 most frequent words of the English language. Many research approaches to readability employ similar ideas to gain insight into the lexical complexity of texts, such as the BIV (basic Italian vocabulary) employed by (Dell’Orletta et al., 2011).

All feature values outside of type-token-ratio (which ignores sentence boundaries) are computed at the sentence level. When a dataset entry comprises multiple sentences, the output for that entry is the average of the values for the sentences that compose it.

3 Results

Table 2 details the average output for each difficulty level and metric. For each metric, the 3 values most

²https://simple.wikipedia.org/wiki/Main_Page

³<https://www.wikipedia.org/>

⁴Color Blindness, The Great Depression, Butterflies, Dogs, Semantics, The Internet, The Moon, Dinosaurs, Economics, Quantum Mechanics’

⁵A1, A2, B1, B2, C1, C2

⁶<https://ai.googleblog.com/2006/08/all-our-n-gram-are-belong-to-you.html>

indicative of low readability are highlighted.

With regards to the selected metrics, GPT-3 does a good job generating texts for the target readability level in the sense that prompts requesting C1 and C2 level material on average yield the least readable texts, whereas lower proficiency levels tend to have scores indicating higher readability.

While no rigorous checks are carried out in this paper in reference to the correlation between these metrics and article readability levels, it’s still possible to examine some of the results to gain some insight into the black box that is GPT-3.

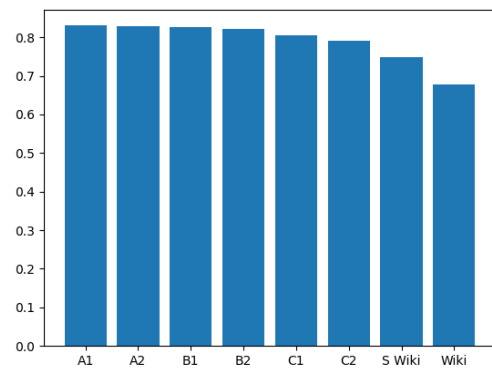


Figure 1: Mean metric values for WIT.

Among the selected measures, WIT (ratio of words in the 3000 most frequent English words) appears to be the most well-behaved, insofar as the prompted CEFR levels mostly maintain their natural progression. Figure 1 shows the isolated scores for the WIT metric, which show a correlation with difficulty. Wikipedia and Simple Wikipedia articles have lower values, however this is possibly connected to those texts being several times longer.




Figure 2: Mean metric values for LLDL.

Level	ASL	NVR	TTR	CPS	LLDL	PTD	WIT
S Wiki	18.145	1.918	0.661	1.629	7.868	4.629	0.749
A1	18.539	1.681	0.657	1.932	8.0	5.113	0.83
A2	20.717	1.477	0.648	2.104	8.843	5.415	0.828
B1	19.124	1.607	0.65	2.114	8.416	5.383	0.826
B2	20.766	1.694	0.632	2.064	9.148	5.64	0.821
C1	21.24	1.848	0.646	1.967	9.147	5.549	0.804
C2	21.94	1.836	0.658	1.94	10.189	5.427	0.791
Wiki	26.131	2.342	0.667	1.876	11.323	5.681	0.678

Table 2: Average metric results for each difficulty level and metric. The 3 highest values are highlighted for each metric. For WIT, the lowest 3 are highlighted instead.

LLDL (length of longest dependency link) is another metric that performs well in terms of level differentiation. It also reflects the common belief that the biggest difficulty spike in language learning - as shown by Figure 2. As with most metrics, however, the natural progression of the CEFR levels is lost, particularly in the middle portions.

4 Discussion and future research

This work attempted to gauge true readability levels of GPT-3 output texts with prompts across 10 topics and CEFR proficiency levels, optimising thus for the second language learning application environment. Outputs computed from the metrics set suggest that GPT-3 is at least somewhat aware of syntax and lexis when performing generation for a target readability level.

Future research into this topic could explore a wider variety of textual metrics, as this work only selected a handful that lent themselves particularly well to interpretation and did not require steep hardware requirements to obtain.

Acknowledgements

References

Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. 2020. [Language models are few-shot learners](#).

Felice Dell’Orletta, Simonetta Montemagni, and Giulia Venturi. 2011. [READ-IT: Assessing readability of Italian texts with a view to text simplification](#).

In Proceedings of the Second Workshop on Speech and Language Processing for Assistive Technologies, pages 73–83, Edinburgh, Scotland, UK. Association for Computational Linguistics.

Edward Loper and Steven Bird. 2002. NLTK: The Natural Language Toolkit. In *Proceedings of the ACL Workshop on Effective Tools and Methodologies for Teaching Natural Language Processing and Computational Linguistics*, pages 62–69. Somerset, NJ: Association for Computational Linguistics. <http://arXiv.org/abs/cs/0205028>.

Christopher D. Manning, Mihai Surdeanu, John Bauer, Jenny Finkel, Steven J. Bethard, and David McClosky. 2014. [The Stanford CoreNLP natural language processing toolkit](#). In *Association for Computational Linguistics (ACL) System Demonstrations*, pages 55–60.

Giulia Venturi, Tommaso Bellandi, Felice Dell’Orletta, and Simonetta Montemagni. 2015. [NLP-based readability assessment of health-related texts: a case study on Italian informed consent forms](#). In *Proceedings of the Sixth International Workshop on Health Text Mining and Information Analysis*, pages 131–141, Lisbon, Portugal. Association for Computational Linguistics.

Resources

All code for this paper is available at https://github.com/waron97/nlp_for_readability_term_paper.